

## Exercise 17 - Investigating Deep Neural Network architectures for cryptocurrency price predictions (in R & Tensorflow)

March 06, 2018

- Installed TensorFlow CPU with R wrapper via Python on Anaconda.
- Installed TensorFlow GPU with CUDA via cuDNN package.

Used raw TensorFlow for predicting ETH (ETH-7, BTC-3):

Performance = 0.54 [single layer]

- Installed TensorFlow CPU with R wrapper via Keras API accessing Python on Anaconda with C++ core.

Used Keras API for TensorFlow for predicting ETH (ETH-7, BTC-3):

Performance = 0.54 (150 it), 0.52 (1500 it. - overfitting) [1 hidden layer with 10 hidden nodes]




Figure 1 - Plots of cost function and accuracy for ETH prediction


Trying by shuffling the input data + printing probabilities:

Performance = 0.53 for 300 iterations [1 hidden layer with 10 hidden nodes]

Trying by shuffling the input data + only taking the 60% confidence or above:

Performance = 0.59 for 100 iterations [1 hidden layer with 10 hidden nodes] (198 preds of 768 in test set)

Using Keras API for TensorFlow GPU (CUDA, cuDNN) via Python on Anaconda (C++ Core) for predicting DASH (DASH Close to DASH-7, BTC-1, BTC-2, BTC-3)




*Figure 2 - Plots of cost function and accuracy for DASH prediction*

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 10)	110
dropout_1 (Dropout)	(None, 10)	0
activation_1 (Activation)	(None, 10)	0
dense_2 (Dense)	(None, 2)	22
activation_2 (Activation)	(None, 2)	0
Total params: 132		
Trainable params: 132		
Non-trainable params: 0		

*Figure 3 - Summary of model*

Performance = 0.52 (1500 it.) [1 hidden layer with 10 hidden nodes] dropout = 0.9

Trying by shuffling the input data into mini-batches + printing probabilities:



*Figure 2 - Plots of cost function and accuracy for DASH prediction*

Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 10)	110
dropout_2 (Dropout)	(None, 10)	0
activation_3 (Activation)	(None, 10)	0
dense_4 (Dense)	(None, 2)	22
activation_4 (Activation)	(None, 2)	0
Total params: 132		
Trainable params: 132		
Non-trainable params: 0		

*Figure 3 - Summary of model*

Performance = 0.57 for 100 iterations [1 hidden layer with 10 hidden nodes]

Trying by shuffling the input data + only taking the 52% confidence or above:

Performance = 0.60 for 100 iterations [1 hidden layer with 10 hidden nodes]  
(216 predictions of 719 in test set)